

# Project 1: Optimizing the Performance of a Pipelined Processor

Zhuohao Li <sup>✉\*</sup>; Xiaoheng Xia <sup>✉†</sup>

April 24, 2022

## 1 Introduction

In this lab, we're gonna rewrite some code in Y86 assemble formats(Part A) and implement new instructions(Part B) by it. Besides, we need to modify the HCL description of the processor to add new instructions, optimize the loop strategy and pipeline to reduce critical path, etc(Part C). It can be simulated and tested by running automated tests.

After this lab, we can master the operation of Y86 related assembly language and the description of Y86 HCL, and have a better understanding of pipeline and program optimization as well.

**Zhuohao** finished Part A,

**Xiaoheng** finished Part B,

Part C and this report.pdf are co-finished by both **Zhuohao** and **Xiaoheng**.

Note that we do not need the GUI interface throughout our experiments, thus it is a good idea to comment out the TK related lines in all Makefile(s).

## 2 Experiments

### 2.1 Part A

In this part, we are required to translate C code into Y86-64 assemble code. This part is not difficult.

#### 2.1.1 Analysis

`sum.y8`

The first question requires writing a function to calculate the sum of sample linked lists. In `example.c` `int sum_list(list_ptr ls)`. We just follow the logic shown in C code. It should consist of some code that sets up the stack

---

\*edith\_lzh@sjtu.edu.cn | 519021911248

†xia.xh@sjtu.edu.cn | 519021910341

structure, invokes a function, and then halts. The most important thing is to figure out the instruction execution flow from C to Y86.

`rsum.y8`

This period is similar to `sum.y8` except change `int sum_list(list_ptr ls)` to `int rsum_list(list_ptr ls)`.

`copy.y8`

It should consist of code that sets up a stack frame, invokes a function `copy` block, and then halts. We need calculate the XOR value of all items in the original array and 3 parameters are passed in when the function is called here.

### 2.1.2 Code

`sum.y8`

---

```

1 # core code for sum_list(list_ptr ls)
2 sumlist:  pushl %ebp      # push stack pointer
3           rrmovl %esp, %ebp
4           xorl  %eax, %eax    #the return val = 0, initialized
5           mrmovl 8(%ebp) , %edx
6           andl  %edx , %edx    #ls ==> %edx, to check if ls is 0/
                                #nullptr or not
7           je   End           #if ls == 0 go to End and return
8 Loop:     mrmovl (%edx) , %ecx    #ls->val ==> %ecx, in loop and
                                #store ls->value
9           addl  %ecx , %eax    #val += ls->val, do adding
10          irmovl $4 , %edi     #store 4
11          addl  %edi , %edx    #next ==> edx, which is equivalent to
                                #ls+4
12          mrmovl (%edx), %esi
13          rrmovl %esi , %edx    #ls->next ==>edx
14          andl  %edx , %edx    #check if ls is 0/nullptr
15          jne Loop            #if ls != 0 go to Loop
16
17 End:      rrmovl %ebp , %esp
18          popl  %ebp
19          ret

```

---

`rsum.y8`

---

```

1 # core code for sum_rlist(list_ptr ls)
2 rsum_list:
3           pushl %ebp        # val is reserved as API register value
4           rrmovl %esp , %ebp
5           pushl %ebx
6           irmovl $4 , %esi
7           subl  %esi , %esp

```

---

```

8    xorl    %eax , %eax  #the return val = 0
9    mrmovl  8(%ebp),%edx
10   andl    %edx , %edx # condition set
11   je End      # if nullptr, go to End and return
12   mrmovl  (%edx) , %ebx #ls->val
13   irmovl  $4 , %esi
14   addl    %esi , %edx
15   mrmovl  (%edx) , %edi
16   rmmovl  %edi , (%esp) #ls->next
17   call    rsum_list    #recurit call
18   addl    %ebx , %eax  #return val+rest_val
19
20 End:  addl    %esi , %esp
21      popl    %ebx
22      popl    %ebp
23      ret

```

---

copy.yo

---

```

1  # core code for copy_block(int *src, int *dest, int len)
2  copy_block:
3      pushl   %ebp
4      rrmovl  %esp, %ebp
5      xorl    %eax , %eax #result = 0, initialized
6
7      mrmovl  12(%ebp) , %edx #edx <==> dest
8      mrmovl  8(%ebp) , %esi  #esi <==> src
9      mrmovl  16(%ebp),%ecx   #ecx <==> len
10     andl    %ecx, %ecx #length
11     je End      # if length <=0, go to End and return
12 Loop: mrmovl (%esi) , %ebx #val = *src
13     rmmovl  %ebx , (%edx)
14     xorl    %ebx , %eax #result ^= val
15     irmovl  $4 , %edi
16     addl    %edi , %edx #src++
17     addl    %edi , %esi #dest++
18     irmovl  $1,%edi
19     subl    %edi , %ecx #length--
20     jne Loop

```

---

### 2.1.3 Evaluation

The evaluation is finished in `./misc`  
`sum.yo` test by `./yas sum.yo;./yis sum.yo`, the result is shown in Fig1.  
which is correct.

```

ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohaoLi/sim/misc$ ./yis sum.yo
Stopped in 45 steps at PC = 0x11. Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%eax: 0x00000000      0x00000cba
%ecx: 0x00000000      0x00000c00
%esp: 0x00000000      0x00000100
%ebp: 0x00000000      0x00000100
%edi: 0x00000000      0x00000004

Changes to memory:
0x00ec: 0x00000000      0x000000f8
0x00f0: 0x00000000      0x0000003d
0x00f4: 0x00000000      0x00000014
0x00f8: 0x00000000      0x00000100
0x00fc: 0x00000000      0x00000011

```

Figure 1: sum.yo test

rsum.yo test by `./yas rsum.yo; ./yis rsum.yo`, the result is shown in Fig2. They're the same.

```

ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohaoLi/sim/misc$ ./yis rsum.yo
Stopped in 85 steps at PC = 0x11. Status 'HLT', CC Z=0 S=0 O=0
Changes to registers:
%eax: 0x00000000      0x00000cba
%esp: 0x00000000      0x00000100
%ebp: 0x00000000      0x00000100
%esi: 0x00000000      0x00000004

```

Figure 2: rsum.yo test

copy.yo test by `./yas copy.yo; ./yis copy.yo`, the result is shown in Fig3. which is correct.

## 2.2 Part B

### 2.2.1 Analysis

In this project, we are going to implement an `iaddl` instruction. The function of `iaddl` is to add a constant value to a register. This requires first using an `irmovl` instruction to set a register to the constant, and then add `addl` instruction to add this value to the destination register.

According to the computing process of `irmovl` and `opl` from the book, we can write the computing process of the `iaddl`.

---

```

1 fetch:
2     icode:ifun <-- M1[PC]
3         rA:rB <-- M1[PC+1]
4         valC <-- M4[PC+2]
5         valP <-- PC+6
6 decode:
7         valB <-- R[rB]

```

```

ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohao1/sim/misc$ ./yis copy.yo
Stopped in 56 steps at PC = 0x11. Status 'HLT', CC Z=0 S=0 O=0
Changes to registers:
%eax: 0x00000000      0x00000cba
%edx: 0x00000000      0x0000002c
%ebx: 0x00000000      0x00000c00
%esp: 0x00000000      0x00000100
%ebp: 0x00000000      0x00000100
%esi: 0x00000000      0x0000000c
%edi: 0x00000000      0x00000001

Changes to memory:
0x0020: 0x00000111      0x0000000a
0x0024: 0x00000222      0x000000b0
0x0028: 0x00000333      0x00000c00

```

Figure 3: copy.yo test

```

8 execute:
9     valE <-- valB + valC
10 # (There's no memory related operation.)
11 write back:
12     R[rB] <-- valE
13 PC update:
14     PC <-- valP

```

In the fetch phase, like most the other instructions, we read the instruction from `M1[PC]` and store it in `icode:ifun`. Then, we read the two operand: `rB`, `V` from `M1[PC+1]` and `M4[PC+2]` respectively. Lastly, we store `PC+6` in `valP`, as to provide a destination for the position of the next instruction.

In the decode phase, we only need to extract `R[rB]` and store it into `valB`, with `valC` already provided as immediate number.

In the execute phase, we use ALU to calculate `valB + valC` and store it into `valE`.

We don't need to read to or write from memory in this instruction.

In the write back phase, we write back the calculated result from `valE` to `R[rB]`

At last, we update Program Counter in the PC update phase.

So that we can modify the `seq-ful.hcl` file accordingly.

### 2.2.2 Code

**Fetch stage:** `IIADDL` should be added in the `instr_valid` section, and it needs register and immediate word `valC`.

```

1 bool instr_valid = icode in
2   { INOP, IHALT, IRRMOVL, IIRMOVL, IRMMOVL, IMRMOVL,
3     IOPL, IJXX, ICALL, IRET, IPUSHL, IPOPL, IIADDL };

```

```

1 # Does fetched instruction require a regid byte?
2 bool need_regids =
3     icode in { IRRMOVL, IOPL, IPUSHL, IPOPL,
4               IIRMOVL, IRMMOVL, IMRMOVL, IIADDL };

```

---

```

1 # Does fetched instruction require a constant word?
2 bool need_valC =
3     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IJXX, ICALL, IIADDL
4               };

```

---

**Decode stage:** Select register B as B source and E destination.

---

```

1 ## What register should be used as the B source?
2 int srcB = [
3     icode in { IOPL, IRMMOVL, IMRMOVL, IIADDL } : rB;
4     icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
5     1 : RNONE; # Don't need register
6 ];
7
8 ## What register should be used as the E destination?
9 int dstE = [
10    icode in { IRRMOVL } && Cnd : rB;
11    icode in { IIRMOVL, IOPL, IIADDL } : rB;
12    icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
13    1 : RNONE; # Don't write any register

```

---

**Execute stage:**

As described before, we select valC and valB as ALU input. The default operation for ALU is add so we do not need to modify that. The condition code should be updated in this instruction.

---

```

1 ## Select input A to ALU
2 int aluA = [
3     icode in { IRRMOVL, IOPL } : valA;
4     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IIADDL } : valC;
5     icode in { ICALL, IPUSHL } : -4;
6     icode in { IRET, IPOPL } : 4;
7     # Other instructions don't need ALU
8 ];
9
10 ## Select input B to ALU
11 int aluB = [
12     icode in { IRMMOVL, IMRMOVL, IOPL, ICALL,
13               IPUSHL, IRET, IPOPL, IIADDL } : valB;
14     icode in { IRRMOVL, IIRMOVL } : 0;

```

**Memory stage and PC update stage:** IADDL operation do not involve in memory operations, and PC update stage do not need to be modified.

It is notable that we should also remove the GUI related code in the `makefile` to avoid compilation errors.

```
./ssim -t ../y86-code/asumi.yo
```

```
(cd ../y86-code; make testssim)
```

```

natsubuntu:~/csharp/project1-handout-2022/sin$seq cd ../j86-code; make tes
tssin)
./seq/ssin +t asm.yo + asm.seq
./seq/ssin +t asm.yo + asm.seq
./seq/ssin +t asm.yo + asm.seq
./seq/ssin +t jcc.yo + jcc.seq
./seq/ssin +t postest.yo + postest.seq
./seq/ssin +t pushquestion.yo + pushquestion.seq
./seq/ssin +t pushstest.yo + pushstest.seq
./seq/ssin +t prog1.yo + prog1.seq
./seq/ssin +t prog2.yo + prog2.seq
./seq/ssin +t prog3.yo + prog3.seq
./seq/ssin +t prog4.yo + prog4.seq
./seq/ssin +t prog5.yo + prog5.seq
./seq/ssin +t prog6.yo + prog6.seq
./seq/ssin +t prog7.yo + prog7.seq
./seq/ssin +t ret-hazard.yo + ret-hazard.seq
grep "ISA Check" * .seq
asm.seq:ISA Check Succeeds
asm.seq:ISA Check Succeeds
cjr.seq:ISA Check Succeeds
jcc.seq:ISA Check Succeeds
postest.seq:ISA Check Succeeds
prog1.seq:ISA Check Succeeds
prog2.seq:ISA Check Succeeds
prog3.seq:ISA Check Succeeds
prog4.seq:ISA Check Succeeds
prog5.seq:ISA Check Succeeds
prog6.seq:ISA Check Succeeds
prog7.seq:ISA Check Succeeds
prog8.seq:ISA Check Succeeds
pushquestion.seq:ISA Check Succeeds
pushstest.seq:ISA Check Succeeds
ret-hazard.seq:ISA Check Succeeds
rm asm.seq asm.seq seq seq seq jcc.seq postest.seq pushquestion.seq pushstest.
seq prog1.seq prog2.seq prog3.seq prog4.seq prog5.seq prog6.seq prog7.seq prog
8.seq ret-hazard.seq
natsubuntu:~/csharp/project1-handout-2022/sin$seq

```

The results are all ISA Check Succeeds.

The results are shown in Figure 5.

7

```
natsu@ubuntu: ~/csapp/project1-handout-2022/sim/seq
natsu@ubuntu:~/csapp/project1-handout-2022/sim/seq$ (cd ../ptest; make SIM=../seq/ssim)
./optest.pl -s ../seq/ssim
Simulating with ../seq/ssim
All 49 ISA Checks Succeed
./jtest.pl -s ../seq/ssim
Simulating with ../seq/ssim
All 64 ISA Checks Succeed
./ctest.pl -s ../seq/ssim
Simulating with ../seq/ssim
All 22 ISA Checks Succeed
./htest.pl -s ../seq/ssim
Simulating with ../seq/ssim
All 600 ISA Checks Succeed
natsu@ubuntu:~/csapp/project1-handout-2022/sim/seq$ (cd ../ptest; make SIM=../seq/ssim TFLAGS=-i)
./optest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 58 ISA Checks Succeed
./jtest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 96 ISA Checks Succeed
./ctest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 22 ISA Checks Succeed
./htest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 756 ISA Checks Succeed
natsu@ubuntu:~/csapp/project1-handout-2022/sim/seq$
```

Figure 5: Regression tests

## 2.3 Part C

Part C is in the `sim/pipe`. We need modify `ncopy.js` and `pipe-full.hcl` so that `ncopy.js` runs as fast as possible.

### 2.3.1 Analysis

The `ncopy` function in Figure 2 copies a `len`-element integer array `src` to a non-overlapping `dst`, returning a count of the number of positive integers contained in `src`. There are a number of ways to optimize this function.

#### 1. Implementing `IIADDL`

We're supposed to add `iaddl` in `pipe-full.hcl`, the implementation of the corresponding `iaddl` is similar to what we did in Part B. By simply substituting for the `iaddl` instruction, we can achieve better CPE.

#### 2. Loop Unrolling



In `ncopy.y`s, we utilize **unrolling**, to cultivate 1 single loop to 4 loops (in the code, `LNpos#`).

It's basic idea is to unwind a loop process by rewriting each loop explicitly in the code. By trading code spaces for time, we can reduce branch penalties, thus optimize the program's execution speed.

Unrolling will increase the parallelism in the system and in VLSI and compiler design, we also utilize that kind of stuff to regularize and reduce our critical path and avoid stalls.

### 3. Bubble elimination

Additionally, there're **hazards** in original code like below:

---

```
1 mrmovl (%ebx), %esi # read val from src...
2 rmmovl %esi, (%ecx) # ...and store it to dst
```

---

Due to RAW hazards, there's a stall to avoid error. By carefully crafting the sequence of instructions, we add another `mrmovl` in the 2 instructions shown above to avoid unnecessary stalls. We can get more bonus from getting prepared value ahead, which is good for loop unrolling.

#### 2.3.2 Code

Though performing loop unrolling could make the code less readable, the code here is rather straightforward. Instead of copying one byte at a time, We copy 4 bytes at a time. If  $len < 4$ , then we do it 1 byte at a time, like the original code.

---

```
1 # You can modify this portion
2 # Loop Header
3     xorl    %eax , %eax # count = 0;
4     iaddl   $-4 , %edx #len = len -4
5     andl    %edx , %edx # len <= 0?
6     jle     remain
7 Loop: mrmovl (%ebx) , %esi # read val from src...
8     mrmovl  4(%ebx),%edi
9     rmmovl  %esi , (%ecx) # ...and store it to dst
10    andl    %esi ,%esi # val <= 0?
11    jle     LNpos1 # if so, goto LNpos1:
12    iaddl   $1 , %eax # count++ using iaddl
13 LNpos1: rmmovl %edi , 4(%ecx)
14    andl    %edi , %edi
15    jle     LNpos2
16    iaddl   $1, %eax
17 LNpos2:mrmovl 8(%ebx) , %esi
18    mrmovl  12(%ebx),%edi
19    rmmovl  %esi ,8 (%ecx)
20    andl    %esi ,%esi
```

```

21     jle LNpos3
22     iaddl $1 , %eax
23 LNpos3: rmmovl %edi , 12(%ecx)
24     andl  %edi , %edi
25     jle   nextLoop
26     iaddl $1, %eax
27 nextLoop:
28     iaddl $16,%ebx
29     iaddl $16,%ecx
30     iaddl $-4,%edx
31     jge Loop
32
33 # maybe just remain less than 3
34 remain: iaddl $4 , %edx # Restore the true len
35     iaddl $-1, %edx
36     jl  Done
37     mrmovl (%ebx) , %esi
38     mrmovl 4(%ebx),%edi
39     rmmovl %esi , (%ecx)
40     andl  %esi ,%esi
41     jle rNpos
42     iaddl $1 , %eax
43 rNpos:
44     iaddl $-1, %edx
45     jl  Done
46     rmmovl %edi , 4(%ecx)
47     andl  %edi , %edi
48     jle  rNpos1
49     iaddl $1, %eax
50 rNpos1:
51     iaddl $-1 , %edx
52     jl  Done
53     mrmovl 8(%ebx) , %esi
54     rmmovl %esi , 8(%ecx)
55     andl  %esi ,%esi
56     jle Done
57     iaddl $1 , %eax

```

---

### 2.3.3 Evaluation

We add the instruction `IADDL` in `pipe-full.hcl` following the same process in Part B.

From Fig 6, we can see that we've successfully implemented `IADDL`.

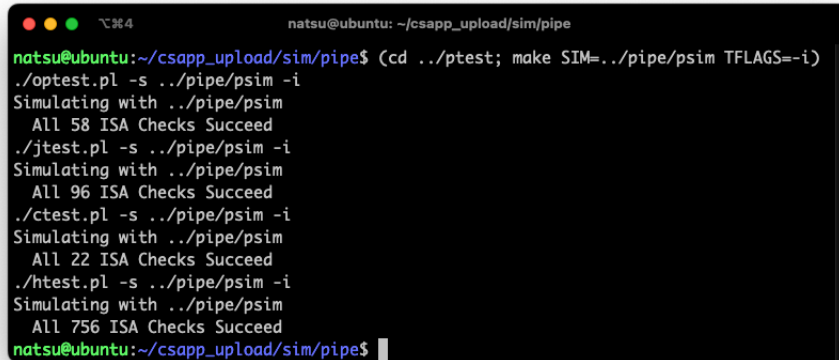
A terminal window with a dark background and light text. The window title is 'natsu@ubuntu: ~/csapp\_upload/sim/pipe'. The prompt is 'natsu@ubuntu:~/csapp\_upload/sim/pipe\$'. The user enters the command '(cd ../ptest; make SIM=../pipe/psim TFLAGS=-i)'. The output shows four test scripts being executed: 'optest.pl', 'jtest.pl', 'ctest.pl', and 'htest.pl'. Each script simulates with './pipe/psim -i' and reports the number of ISA checks that succeeded: 58, 96, 22, and 756 respectively. The prompt returns to 'natsu@ubuntu:~/csapp\_upload/sim/pipe\$'.

Figure 6: IADDL implementation check

Following the guidance in assignment. We evaluate the correctness and efficiency.

#### Correctness

Using commands below to test the correctness.

```
ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohaoli/sim/pipe$ ./correctness.pl
Simulating with instruction set simulator yis
ncopy
```

```
ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohaoli/sim/pipe$ ./correctness.pl -p
Simulating with pipeline simulator psim
ncopy
```

And the result is all CORRECT!(Fig 7)

```
256    OK
68/68 pass correctness test
```

Figure 7: correctness check

efficiency Using ./benchmark.pl in ./pipe

```
ed5@ubuntu:~/Desktop/Arch/arch_lab1_zhuohaoli/sim/pipe$ ./benchmark.pl
ncopy
```

And the Average CPE is **9.48**!(Fig. 8)

50	369	7.38
51	376	7.37
52	376	7.23
53	385	7.26
54	395	7.31
55	402	7.31
56	402	7.18
57	411	7.21
58	421	7.26
59	428	7.25
60	428	7.13
61	437	7.16
62	447	7.21
63	454	7.21
64	454	7.09
Average	CPE	9.48
Score	60.0/60.0	

Figure 8: speed upon benchmark

Indicating that we have finished this lab.

In addition, I've modified `benchmark.pl` to test our solution on a larger scale. We achieved a average CPE of 7.47 when the upper bound is set to 256, as shown in Fig 9

246	1643	6.68
247	1650	6.68
248	1650	6.65
249	1659	6.66
250	1669	6.68
251	1676	6.68
252	1676	6.65
253	1685	6.66
254	1695	6.67
255	1702	6.67
256	1702	6.65
Average	CPE	7.47
Score	60.0/60.0	

Figure 9: Larger scale benchmark

### 3 Conclusion

From this lab, we've familiarize ourselves with the Y86 instruction set. Now, we are able to write an assembly program from scratch, implement a new instruction by analyzing its computing process, and optimize a given assembly function by carefully avoiding hazards to achieve best CPE. Here is a summary of the completion of the three parts:

- Part A:
  - + We familiarized ourselves with the Y86 instruction set.
  - + We translated a C program into a Y86 assembly program.
- Part B:
  - + We wrote out the computing process of the instruction `IADDL`
  - + We modified the file `seq-full.hcl` to implement the instruction `IADDL`
- Part C:
  - + We substitute for the `IADDL` instruction to speed up the process.
  - + We did a 4-way loop unrolling to reduce branch penalties and optimize branch prediction.
  - + We reordered the instructions to avoid RAW hazards and eliminate bubbles.

### 3.1 Feelings

This lab allows both of us to gain a deeper understanding in concepts and optimizations in Computer Architecture. By various hands-on approach, we debugged the problems and optimized the `ncopy` function step by step. We are grateful for this opportunity of being introduced to this valuable lab, and we would like to thank Prof. Shen and the TAs for their support.