

Fundamentals

≡ Author	Dawson Lee
≡ Last Update	2022/03/07

Introduction

Evolution of Processors

x86 Manufacturers

ISA

Die

First Microprocessor

General

Single Processor Performance

Effects

Current Trends

Classes of Computers

Flynn's Taxonomy

Two Kinds of Parallelism in Applications

Four Major Ways for Exploiting Parallelism

Define Computer Architecture

Contents of ISA

Addressing Mode

MIPS ISA

Addressing Mode Examples

Micro-architecture

Concepts

Computer Architecture in General

Trends in Technology

Five Critical Implementation Techs

IC technology

Quantify

Feature size

Trends in Power and Energy in ICs

Power and Energy

Thermal Design Power (TDP)

Clock Rate

Dynamic Energy and Power

Power

Techniques

Static Power

Trends in Cost

Dependability

Measuring Performance

Typical Performance Metrics

Benchmarks

Quantitative Principles

Principles for Computer Design

Amdahl's Law

Introduction

Evolution of Processors

- 8086: x86 Architecture
- Pentium
- IntelCore

x86 Manufacturers

- Intel
- AMD
- VIA
- Others are acquired or discontinued

ISA

- x86 is CISC (complex instruction set computer)
 - x86 has varying length instructions
- RISC (reduced instruction set computer)

Die

Wafer consists of many Dies.

First Microprocessor

- Intel 4004
- Max CPU clock rate: 108 kHz

General

- Transistors doubles every 1.5-2.0 yrs
- Process speed doubles every 1.5-2.0 yrs
- DRAM size doubles every 1.5-2.0 yrs

Single Processor Performance

- 2003: from uni-processor to multi-processor
- Pipeline (ILP) reaches its limit
- Power limit
- Move to multi-processor
- Clock rate 1% growth

Effects

- enhanced capability
- lead to new class of computers
- dominance of microprocessor-based computers
- software development shift focus on productivity

Current Trends

- DLP
- TLP
- RLP

Classes of Computers

▼ class

1. Personal Mobile Device
2. Desktop Computing
3. Servers
4. Clusters / Warehouse Scale Computers
5. Supercomputers (require faster network than WSCs)
6. Embedded Computers

(refer to textbooks for details)

Flynn's Taxonomy

- Single instruction stream, single data stream (SISD)
- SIMD (GPU, AVX extensions, vector arch)
- MISD (no commercial implementation)
- MIMD (tightly-coupled / loosely-coupled)

Two Kinds of Parallelism in Applications

- Data-Level Parallelism (operate on many data at the same time)
- Task-Level Parallelism (tasks can operate independently)
- example: web crawler
 - crawl web pages
 - parse HTML data (data-level)
 - run parse HTML task (task-level)

Four Major Ways for Exploiting Parallelism

- ILP (pipelining, data-level)
- Vector arch / GPU (data-level)
- Thread-Level Parallelism (multi-core, data-level or task-level)
- Request-Level Parallelism (clusters, data-level or task-level)

Define Computer Architecture

- old view: ISA
- ISA is same as programming model
- abstracts hardware and software interface

Contents of ISA

- register (size, number)
- class of ISA (register-memory, complex or load/store)
- memory addressing (byte addressing, little/big-endian)
- addressing modes
- instruction operands (e.g. RISC-V 3 operands, x86 only 2)
- available operations (e.g. RV32I doesn't support hardware mul and div)
- control flow instructions (e.g. ret or jalr in RISC-V reads return address from ra, but x86 reads from stack and pop)
- instruction encoding (fixed length / variable length)
- etc.

Addressing Mode

- register, $r4 \leftarrow r4 + r3$
- immediate, $r4 \leftarrow r4 + 3$
- displacement, $r4 \leftarrow r4 + M[100+r1]$
- register deferred, $r4 \leftarrow r4 + M[r1]$
- etc. refer to textbook

MIPS ISA

- `| op 5 | rs 5 | rt 5 | rd 5 | shamt 5 | funct 6 |` (32 bits)

op: operation
rs: register first source
rt: register second source
rd: register destination
shamt: shift amount
funct: function code

Addressing Mode Examples

register addressing (no shamt)

immediate addressing (rd-funct as operand)

base (displacement) addressing (rd-funct as offset + base register)

pc-relative addressing (rd-funct as offset + PC)

Micro-architecture

- micro-arch, also called computer organization, is the way a given ISA is implemented on a processor.
- a given ISA can be implemented with different micro-archs

Concepts

- pipelining
- hierarchical memory organization
- cache
- cache coherence
- branch prediction
- super-scalar
- out-of-order execution
- register renaming
- multi-processing and multi-threading

Computer Architecture in General

- working in constraints
- market target?
- cost/performance?
- tradeoff in material and process
- Computer Architecture is about designing the organization and hardware to meet goals
- and functional requirements

Trends in Technology

Five Critical Implementation Techs

1. IC technology
2. semiconductor DRAM
3. semiconductor flash
4. magnetic disk technology
5. network technology

IC technology

vacuum tube → transistor → semiconductor (gates, memory cells, interconnections)

Quantify

- bandwidth or throughput: total work done in a given time
- latency or response time: time between start and completion of an event
- example: in a pipelined processor, it's running a series of add instruction. There might be multiple add running at the same time.
 - throughput: how many add can be issued in a given time
 - latency: for one add instruction, how long does it take from issue to complete.

- bandwidth outpaced latency

Feature size

- transistor size in x or y dimension
- transistor density increases
- wire latency do not increase that fast

Trends in Power and Energy in ICs

Power and Energy

power is unit time energy

Thermal Design Power (TDP)

- characterize sustained power consumption
- use as target for power supply and cooling
- lower than peak power, higher than average power consumption

Clock Rate

can be reduced dynamically to reduce power consumption

Dynamic Energy and Power

- dynamic energy per transistor
 - used for a transistor switching from 0 to 1 / 1 to 0
 - $\frac{1}{2} \text{capacitive load} \times \text{voltage}^2$
 - related to number of transistor and technology
- dynamic power per transistor
 - $\frac{1}{2} \text{capacitive load} \times \text{voltage}^2 \times \text{frequency switched}$
- voltage is the key
 - voltage of processor has become lower

Power

- 130W, maximum for air cooling
- we use energy of a specific task to compare CPUs

Techniques

- turn off clock
- dynamic voltage-frequency scaling
- low power state for DRAM, disks
- overclocking, turning off cores

Static Power

- current \times voltage
- scales with number of transistors
- power gating: turning off the power supply

Trends in Cost

- refer to slides and textbook
- I don't want to type exactly the same thing in slides
- From my perspective, quantifying these things is the core of the book CAAQA. However, this is not our course orients to.

Dependability

- $MTTF \sim \text{Exp}(n)$
- $MTTF \text{ Components} \sim \text{Exp}(n_1 + n_2 + \dots)$
- $E(\text{Exp}(x)) = \frac{1}{x}$
- Serial MTTF

$$\frac{1}{\frac{1}{MTTF} + \frac{1}{MTTF}}$$

- Redundant MTTF

$$\frac{MTTF^2}{2 \times MTTR}$$

Measuring Performance

Typical Performance Metrics

- latency (response time)
- throughput (bandwidth)

Benchmarks

- a common program for testing the execution times of computers
- e.g. kernels (matrix multiply), top programs (e.g. sorting), synthetic benchmarks
- we use benchmark suite (SPEC)
- SPEC uses performance ratio
- and Geometric Mean

Quantitative Principles

Principles for Computer Design

take advantage of parallelism
principle of locality
focus on common case

Amdahl's Law

- performance improvement of using a new feature is limited by the fraction of the time the new feature can be used
- formula

$$\text{Execution time}_{\text{new}} = \text{Execution time}_{\text{old}} \times \left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)$$

$$\text{Speedup}_{\text{overall}} = \frac{\text{Execution time}_{\text{old}}}{\text{Execution time}_{\text{new}}} = \frac{1}{(1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}}}$$

- specific feature can speedup the program. **fraction is the part that actually accelerated (10%...).** **Speedup is the times that the acceleration (2 times ...)**
- can compute fraction, speedup.